

Data Structures

I didn't need to add or change any data structures, I did and the abstract methods `void ensure_max_priority(void);` and `bool compare_priority(const struct list_elem *a, const struct list_elem *b);` in thread.h, and `bool compare_sema_priority(const struct list_elem *a, const struct list_elem *b);` in synch.h, these methods will be detailed in the algorithms section.

Algorithms

My solution for part A is based around the `ensure_max_priority` and `compare_priority` methods added to thread.c.

`ensure_max_priority` begins by checking if there are threads on the ready list, if there are no threads the method can exit saving execution time. The method then retrieves the top of the ready list and compares its priority to the current thread, if the current thread has lower priority then it yields the processor. The ready list is assumed to be ordered by priority so only the top entry needs to be checked.

`ensure_max_priority` is used after a new thread is created and after a thread's priority is changed. `thread_set_priority` will exit before making changes and calling `ensure_max_priority` if the new priority equals the current priority.

The ready list is kept ordered because threads are added to it with the `list_insert_ordered` method from list.c which runs in $O(n)$ as an average case. The ordering is determined by the `compare_priority` method.

`compare_priority` is a simple method that takes two list elements, extracts the threads then compares their priority. `compare_priority` returns true if the first element has a higher priority than the second otherwise it returns false.

My solution for part B reuses the `ensure_max_priority` and `compare_priority` methods and also adds the `compare_sema_priority` method in synch.c.

`compare_sema_priority` is used to sort the list of waiting semaphores for a conditional variable, sorting is done in the `cond_signal` method with `list_sort` from list.c which runs in $O(n \lg n)$ time.

`compare_sema_priority` takes two semaphores and compares their priority returning true if the first has higher priority and false otherwise.

`compare_sema_priority` is much like `compare_priority` except first it checks that there are threads waiting on both semaphores. If b has no waiting threads then the method can return immediately, either false if a also has nothing waiting or true if it does. Likewise if a has no threads but b does then false can be returned immediately. These checks save processing time on sorting the waiting threads. The waiting threads are sorted using `list_sort` with `compare_priority` from part A. Once sorted the top threads can have priorities compared as in `compare_priority`.

In `sema_up` waiting threads are sorted using `list_sort` with `compare_priority` before the top waiting thread is unblocked. `ensure_max_priority` is used at the end of the method to ensure that the unblocked thread will gain the cpu if it has higher priority than the current thread.

Synchronization

To avoid race conditions `ensure_max_priority` disables interrupts, as do `thread_set_priority` and `thread_get_priority`.