

# Database Concurrency Control

SWEN304/SWEN439

Lecturer: Dr Hui Ma

**Engineering and Computer Science**



Slides by: Pavle Morgan & Hui Ma

# Outline

- Transaction schedules
- Basic locks and basic locking rules
- Lock conversion
- Lost update and locking
- Protocols to insure isolation property of concurrent transactions
- Dead lock and dead lock prevention protocols
- Starvation
- Phantom record
  - *Readings from the textbook:*
    - *Chapter 21: Section 21.5,*
    - *Chapter 22 : Sections 22.1, 22.2, and 22.5*
  - *PostgreSQL Manulas*

# Database Concurrency Control

- **Purpose of Concurrency Control**
  - To enforce Isolation (through mutual exclusion) among conflicting transactions
  - To preserve database consistency through consistency preserving execution of transactions
  - To resolve read-write and write-write conflicts
- Example: In concurrent execution environment if  $T_1$  conflicts with  $T_2$  over a data item  $A$ , then the existing concurrency control decides if  $T_1$  or  $T_2$  should get the  $A$  and if the other transaction is rolled-back or waits

# Transaction Schedules

- **Transaction schedule or history:** When transactions are executing concurrently in an interleaved fashion, the order of execution of operations from the various transactions forms what is known as a **transaction schedule (or history)**
- A **schedule (or history)**  $S$  of  $n$  transactions  $T_1, T_2, \dots, T_n$  is an ordering of the operations of the transactions subject to the constraint that, for each transaction  $T_i$  that participates in  $S$ , the operations of  $T_i$  in  $S$  must appear in the same order in which they occur in  $T_i$
- Note, however, that operations from other transactions  $T_j$  **can be interleaved** with the operations of  $T_i$  in  $S$

# Transaction Schedules based on Serializability

- **Serial schedule:** A schedule  $S$  is **serial** if, for every transaction  $T$  participating in the schedule, all the operations of  $T$  are executed consecutively in the schedule. Otherwise, the schedule is called **nonserial schedule**
- **Serializable schedule:** A schedule  $S$  is **serializable** if it is equivalent to some serial schedule of the same  $n$  transactions
- **Result equivalent:** Two schedules are called *result equivalent* if they produce the same final state of the database

# Schedules based on Serializability (3)

- Being serializable is not the same as being serial
- Being serializable implies that the schedule is a **correct** schedule
  - It will leave the database in a consistent state
  - The interleaving is appropriate and will result in a state as if the transactions were serially executed
    - will achieve efficiency due to concurrent execution

# Transaction Schedules based on Serializability

## Practical approach:

- Come up with methods (protocols) to ensure serializability
- It's not possible to determine when a schedule begins and when it ends
- Hence, we reduce the problem of checking the whole schedule, to checking only a *committed project* of the schedule (i.e. operations from only the committed transactions.)
- Current approach used in most DBMSs:
  - Use of locks with two-phase locking

# Locking

- Locking is the most frequent technique used to control concurrent execution of database transactions
- Operating systems provide a binary locking system (*lock* and *unlock*) that is too restrictive for database transactions
- That's why DBMS contains its own lock manager
- A  $\text{lock\_value}(X)$  is variable associated with (each) database data item  $X$
- The  $\text{lock\_value}(X)$  describes the status of the data item  $X$ , by telling which operations can be applied to  $X$



# Kinds of Locks

- Generally, the lock manager of a DBMS offers two kinds of locks:
  - shared (read) lock
  - exclusive (write) lock
- If a transaction  $T$  issues a  $\text{read\_lock}(X)$  command, it will be added to the list of transactions that share lock on item  $X$ , unless there is a transaction already holding write lock on  $X$
- If a transaction  $T$  issues a  $\text{write\_lock}(X)$  command, it will be granted an exclusive lock on  $X$ , unless another transaction is already holding lock on  $X$
- Accordingly,  
 $\text{lock\_value} \in \{\text{read\_lock}, \text{write\_lock}, \text{unlocked}\}$

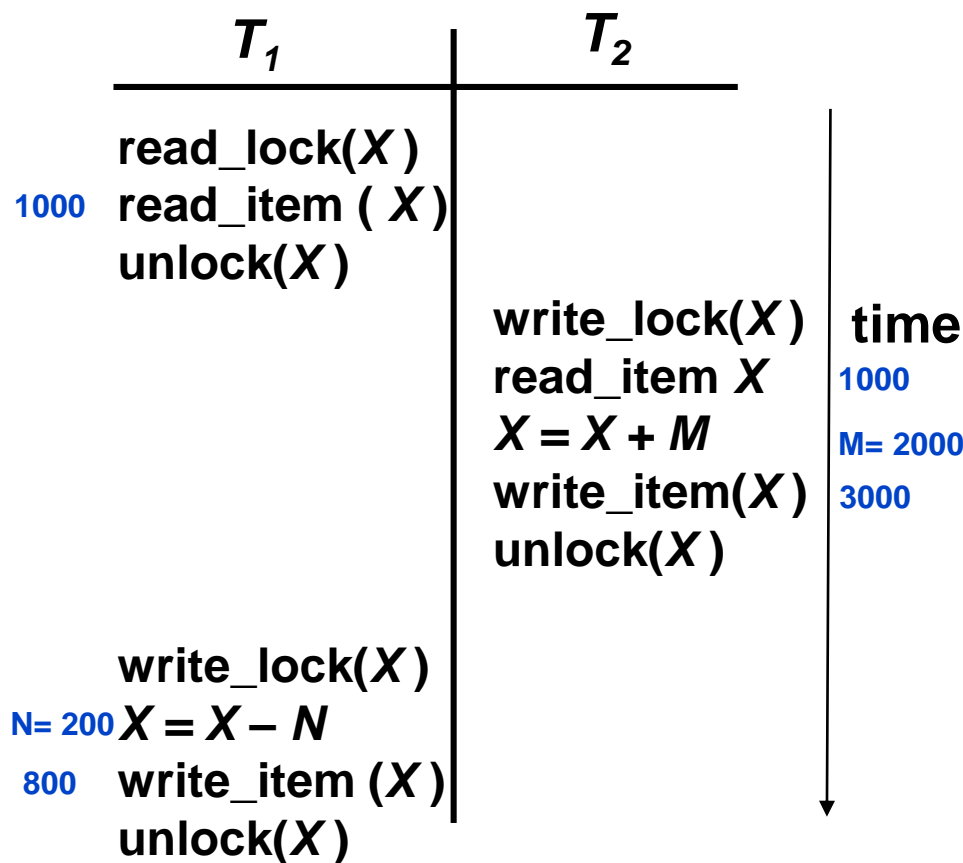
# Basic Locking Rules

- The basic locking rules are:
  - $T$  must issue a **read\_lock( $X$ )** or **write\_lock( $X$ )** command before any **read\_item( $X$ )** operation
  - $T$  must issue a **write\_lock( $X$ )** command before any **write\_item( $X$ )** operation
  - $T$  must issue an **unlock( $X$ )** command when all **read\_item( $X$ )** or **write\_item( $X$ )** operations are completed
- Some DBMS lock managers perform automatic locking by granting an appropriate database item lock to a transaction when it attempts to read or write an item into database
- So, an item lock request can be either explicit, or implicit

# Lock Conversion

- A transaction  $T$  that already holds a lock on item  $X$  can convert it to another state:
  - $T$  can **upgrade** a **read\_lock( $X$ )** to a **write\_lock( $X$ )** if it is the only one that holds a lock on the item  $X$  (otherwise,  $T$  has to **wait**)
  - $T$  can always **downgrade** a **write\_lock( $X$ )** to a **read\_lock( $X$ )**

# Lost Update Problem and Locking



- The problem is that  $T_1$  releases lock on  $X$  too early, allowing  $T_2$  to start updating  $X$
- We need a protocol that will guarantee database consistency

# 2-Phase Locking Techniques: The algorithm

- **Two Phases**

- **(a) Locking (Growing) Phase:** A transaction applies locks (read or write) on desired data items one at a time.
- **(b) Unlocking (Shrinking) Phase:** A transaction unlocks its locked data items one at a time.
- **Requirement:** For a transaction these two phases must be mutually exclusively, that is, during locking phase unlocking phase must not start and during unlocking phase locking phase must not begin.

# Strict 2-Phase Locking

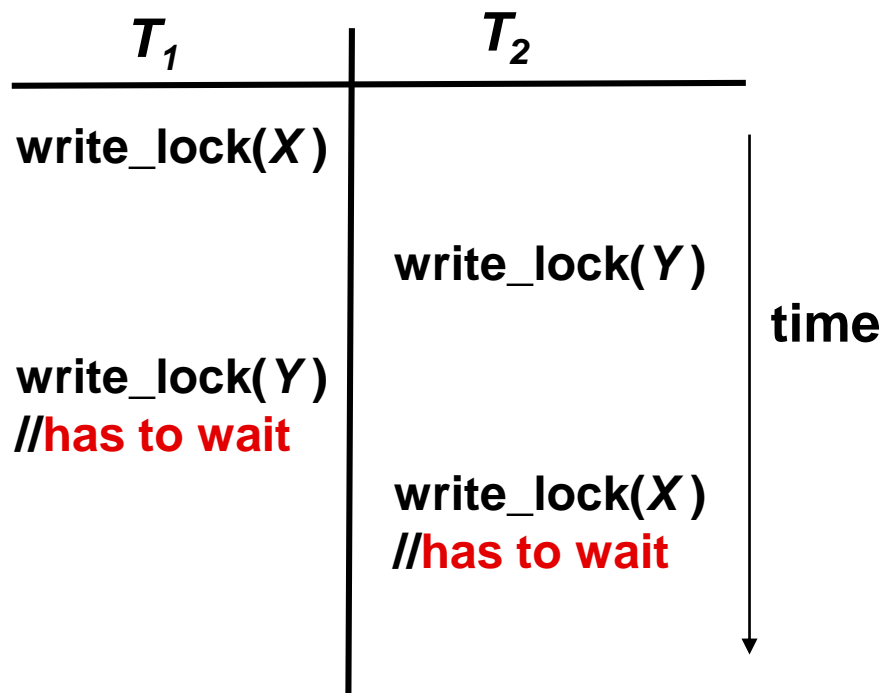
- Protocol:
  - All lock operations of a transaction  $T$  **must precede** the first unlock operation
  - A transaction  $T$  does not release any of exclusive locks until after it **commits** or **aborts**
- Comments:
  - No other transaction can read or write an item  $X$  that is written by  $T$  unless  $T$  has **committed**
  - The strict 2-phase locking protocol is **safe** for all transaction anomalies mentioned so far
  - It is also called **read committed** protocol, because transactions are allowed to read only committed database items

# Undesirable Effects of Locking

- 2-phase locking can introduce some undesirable effects:
  - waits,
  - deadlocks,
  - Starvation
- **Waits** relate to the fact that a transaction wanting to acquire a lock on a database item  $X$  has to wait if another transaction has already acquired an exclusive lock on  $X$

# Deadlock

- **Deadlock** is also called **deadly embrace**
- Typical sequence of operations is given in the following diagram



- $T_1$  acquired exclusive lock on  $X$
- $T_2$  acquired exclusive lock on  $Y$
- No one can **finish**, because **both** are in the **waiting** state



# Deadlock

- Deadlock occurs when:
  - Each transaction  $T_i$  in a set of *two or more transactions*  $T = \{T_1, T_2, \dots, T_n\}$  is **waiting** for some item  $X$  that is **locked** by some other transaction  $T_j$
- In other words:
  - A number of transactions (greater than one) hold lock on one item and wait to acquire another
  - **None** of the waiting transactions can acquire locks on all necessary items


# Deadlock Examples

a)

- $T_1$  has locked  $X$  and waits to lock  $Y$
- $T_2$  has locked  $Y$  and waits to lock  $Z$
- $T_3$  has locked  $Z$  and waits to lock  $X$

b)

- Both  $T_1$  and  $T_2$  have acquired sharable locks on  $X$  and wait to lock  $X$  exclusively
- A dead-lock may be represented using a cyclic **wait-for graph**



# Deadlock Prevention Techniques (1)

- We distinguish between deadlock **prevention** and deadlock **detection** techniques
- Deadlock prevention techniques:
  - **Conservative 2-phase lock** protocol: lock all items in advance, if any of them cannot be obtained, none of the item are locked; try again later
  - *Timestamp techniques:*
    - **Wait–Die** protocol: if  $TS(T_i) < TS(T_j)$  ( $T_i$  is older than  $T_j$ ) then  $T_i$  is allowed to wait. Otherwise ( $T_i$  is younger than  $T_j$ ) abort  $T_i$  (dies) and restart it later with the same timestamp
    - **Wound–Wait** protocol: if  $TS(T_i) < TS(T_j)$ , ( $T_i$  is older than  $T_j$ ) then abort  $T_j$  ( $T_i$  wounds  $T_j$ ) and restart it later with the same timestamp. Otherwise ( $T_i$  is younger than  $T_j$ )  $T_i$  is allowed to wait

## Deadlock Prevention Techniques (2)

- **No Waiting (NW)** protocol: if unable to get a lock, immediately abort and restart again after a certain time
- **Cautious Waiting (CW)** protocol: if  $T_j$  is not blocked, then  $T_i$  is blocked and allowed to wait; otherwise abort  $T_i$

# Conservative 2-Phase Locking Protocol

- **Conservative 2-Phase Locking Protocol:**
  - A transaction has to lock all items it will access **before** it begins to execute
  - If it cannot acquire any of its locks, it releases all items, **aborts**, and tries again,
- Comments:
  - Deadlock can't occur because no hold-and-wait
  - Once it starts, a transaction can only release its locks
- Problems:
  - What if a transaction cannot predetermine all items it is going to use? (e.g. a sequence of interactive SQL statements comprising one database transaction)
  - What if a database item that is already locked by another transaction will be released very soon? (i.e. the transaction is aborted in vane)

# Deadlock Detection Schemes

- Deadlock **prevention** is justified if transactions are long and use many items, or transaction load is very heavy
- In many practical situations it is advantages to do no deadlock prevention but to **detect dead locks** and then abort at least one of the transactions involved
- Deadlock **detection** schemes are:
  - Deadlock detection using **wait-for graph**
  - **Timeouts protocol**

# Deadlock Detection Protocols

- Deadlock detection using a **wait-for graph**:
  - Construct a wait-for graph where each transaction has its node
  - If  $T_i$  waits on  $T_j$ , construct a directed edge from  $T_i$  to  $T_j$
  - If there is a cycle detected, select a 'victim' and abort it
  - Victim selecting algorithm should select and abort transactions that made the least number of updates
- **Timeouts** protocol:
  - If a transaction waits longer than a specified amount of time, it gets aborted
  - Here, deadlock is only supposed, not proved

# Starvation

- **Starvation** occurs when a transaction can not make any progress for an indefinite period of time, while other transactions proceed
  - can occur when waiting protocol for locked items is **unfair** (used stacks instead of queues)
  - In a deadlock resolution it is possible that the same transaction may consistently be selected as victim and rolled-back
  - This limitation is inherent in all priority based scheduling mechanisms
- Wound-Wait and Wait-Die schemes can avoid starvation, because the aborted transactions restart with the same original timestamp



# Granularity of Items

- Until now, we used the term `data item' without specifying its exact meaning
- In the context of the concurrency control, a data item can be:
  - A **field** of a database record,
  - A database **record**,
  - A disk **block**,
  - A whole **file**,
  - A whole **database**
- The **coarser** data item granularity is, the more contention between transactions will occur, and less productive the DBMS will be (more waits or aborts)

# Granularity of Items (continued)

- The **finer** data granularity, the higher locking overhead of the DBMS lock manager (due to many locks and unlocks)
- The best item size depends on the type of a transaction:
  - If a transaction accesses a small number of records, than
 

data item = record
  - If a transaction accesses a large number of records in the same file, then
 

data item = file
  - Some DBMS automatically change granularity level with regard to the number of records a transaction is accessing (attempting to lock)

# Phantom Record

- A transaction **locks** database items that satisfy certain selection condition and updates them
- During that update, another transaction inserts a **new** item that satisfies the same selection condition
- After the update, but inside the same transaction, we suddenly discover the existence of a database item that has not been updated although it should have been (since it satisfies the selection condition)
- This database item, called a “**phantom record**”, appeared because it did not exist when locking has been done

# Summary

- Basic locks:
  - Shareable,
  - Exclusive
- To avoid all update anomalies:
  - Lost Update,
  - Unrepeatable Read, and
  - Dirty Read

locks should be released only just after the COMMIT point

- Two phase locking protocol may introduce:
  - Waits,
  - Deadlocks,
  - Starvation

# Summary (continued)

- There are many deadlock prevention schemes, but no one is ideal
- In the context of the concurrency control, a database item can be:
  - a field of a database record (tuple),
  - a database record,
  - a disk block,
  - a whole table
  - a whole file,
  - a whole database
- Each data item granularity has advantages and disadvantages, but database record granularity is desirable
- Phantom record may appear if a finer granularity than a table is used