

# Recovery

SWEN304/SWEN439

Lecturer: Dr Hui Ma

**Engineering and Computer Science**



Slides by: Pavle Morgan & Hui Ma

# Plan For The Recovery Topic

- Transaction Log File
- Classification of database recovery procedures
  - Deferred database update
  - Immediate database update
- Checkpoints
  - *Readings:*
    - *Chapter 23 of the textbook*

# Purpose of Database Recovery

- To bring the database into the last consistent state, which existed prior to the failure
- To preserve transaction properties (**A**tomicity, **C**onsistency, **I**solation and **D**urability)
- Example:
  - If the system crashes before a fund transfer transaction completes its execution, then either one or both accounts may have incorrect value.
  - Thus, the database must be restored to the state before the transaction modified any of the accounts

# Transaction Log

- To be able to recover from failures DBMS maintains a log file
- Recovery from any type of failure requires:
  - data values prior to modification (BFIM - BeFore IMage)
  - new values after modification (AFIM – AFter IMage)
- Typically, a log file contains records with the following contents:

[start\_transaction,  $T$ ] (\*  $T$  is a transaction id\*)

[write\_item,  $T$ ,  $X$ , old\_value, new\_value]

[read\_item,  $T$ ,  $X$ ] (\*optional\*)

[commit,  $T$ ]

[abort,  $T$ ]

## Transaction Roll-back (Undo) and Roll-Forward (Redo)

- To maintain atomicity, a transaction's operations are **redone** or **undone**
  - **Undo**: Restore all BFIMs on to disk (Remove all AFIMs)
  - **Redo**: Restore all AFIMs on to disk
- Database recovery is achieved either by performing only Undos or only Redos or by a combination of the two
- These operations are recorded in the log as they happen

# Classification of database recovery procedures

- According to the type of a failure, recovery procedures are classified as:
  - Recovery from a **catastrophic** failure (like disk crash), and
  - Recovery from a **noncatastrophic** failure
- Recovery from a **catastrophic** failure is based on restoring a database **back\_up** copy by redoing operations of committed transactions (stored in archived **log** files) up to the time of the failure

# Noncatastrophic Failures

- A computer failure (system crash):
  - E.g. a hardware, software or network error occurs in the computer system
- A transaction or system error:
  - E.g. integer overflow, division by zero, logical error, user interruption
- Local errors or exception conditions detected by transactions
  - E.g. Data not found, exception condition
- Concurrency control enforcement:
  - E.g. violate serializability, deadlock

# Classification (continued)

- If a database becomes inconsistent due to a **noncatastrophic** failure, the strategy is to reverse only those changes that made database inconsistent
- It is accomplished by **undoing** (and sometimes also **redoing**) some operations, with the use of an in memory log file
- From now on we consider only recovery from non-disk crash failures (we suppose data on disk are safe)
- The recovery from noncatastrophic failures can be based on many algorithms, as:
  - **Deferred** update,
  - **Immediate** update, and
  - Shadow update (not discussed)



# Database Recovery

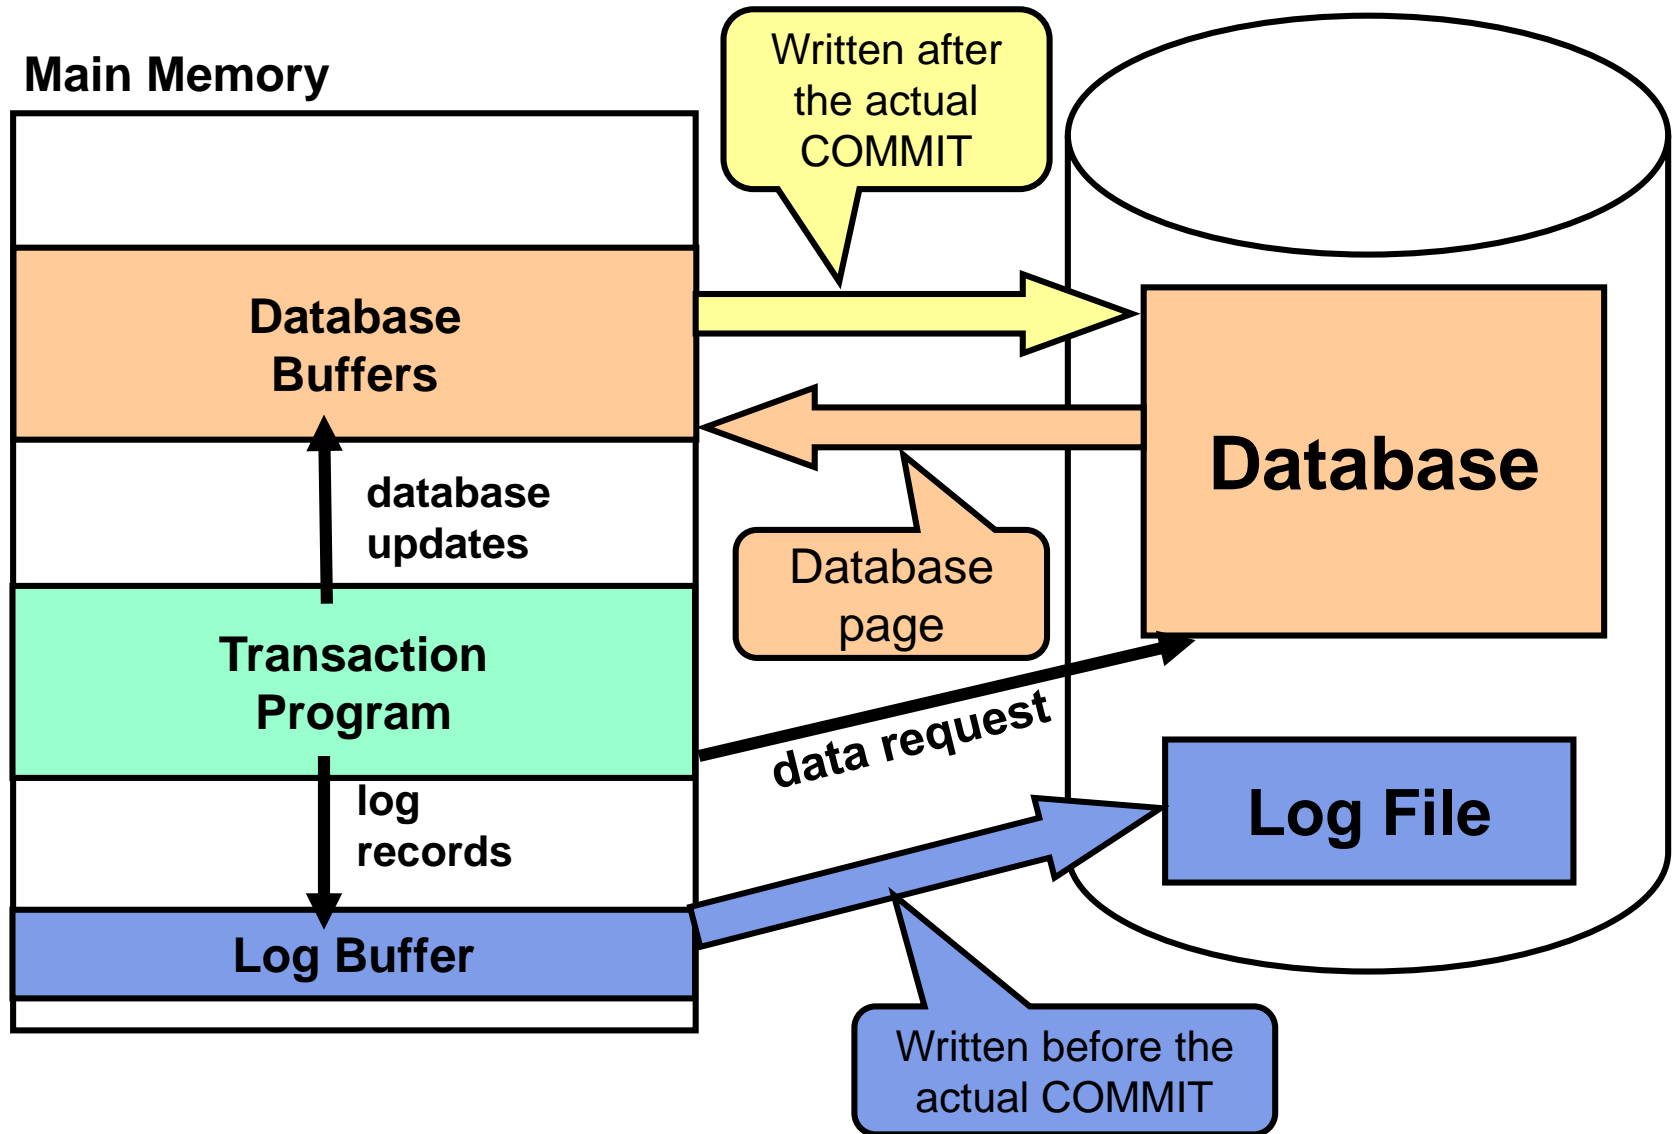
## ■ Data Update

- **Deferred Update:** All modified data items in the cache are written either after a transaction ends its execution or after a fixed number of transactions have completed their execution
- **Immediate Update:** As soon as a data item is modified in cache, the disk copy is updated
- **Shadow update:** The modified version of a data item does not overwrite its disk copy but is written at a separate disk location

# Deferred Update

- The idea:
  - Postpone updates to the database until a transaction reaches its **commit** command
- Updates are recorded in a log file and in cache buffers with database pages (all in RAM)
- When the COMMIT is reached, before it is executed all log updates are first force written to the **log file** on disk, and then the transaction commits
- After that, corresponding updates are written from **buffers** to the database

# Deferred Update Layout



# Deferred Update (continued)

- If a transaction fails before reaching COMMIT, there is no need to make any recovery
- If a system crash occurs after COMMIT, but before all changes are recorded in the database on disk, **Redo** of operations is needed,
  - The operations have to be redone from the log file (that is already on disk) to the database
  - Using after images, AFIMs (item values intended to be written to the database) to perform redo
  - Deferred update recovery log file has to contain only after images - the **new** database item values

# Deferred Update (An Example)

Operations of $T_1$	Generated log records
start $T_1$	begin, $\langle T_1 \rangle$
read_item( $A$ ) // $A = 4000$	
$A = A - 1000$	
write_item( $A$ )	$\langle T_1, \text{write\_item}(A), 3000 \rangle$
read_item( $B$ ) // $B = 0$	
$B = B + 1000$	
write_item( $B$ )	$\langle T_1, \text{write\_item}(B), 1000 \rangle$
read_item( $A$ ) // $A = 3000$	
$A = A - 1000$	
write_item( $A$ )	$\langle T_1, \text{write\_item}(A), 2000 \rangle$
read_item( $C$ ) // $C = 0$	
$C = C + 1000$	
write_item( $C$ )	$\langle T_1, \text{write\_item}(C), 1000 \rangle$
commit $T_1$	
	//force write to the log on disk
	commit, $\langle T_1 \rangle$ // <b>actual</b>


# Deferred Update

- To finish a transaction, DBMS has force stored log records on disk
- If the system fails after that point, DBMS will use the log records to REDO changes in the database
- It is sufficient to redo only the **last** written value (after image) of every item changed
- So, redoing starts from the end of the log file and maintains a list of already redone database items

# Checkpoints

- Some DBMS use **CHECKPOINT** records in the log file to prevent unnecessary redo operations
- To take a checkpoint record, DBMS has:
  - To suspend temporarily operations of all transactions,
  - To force write results of all update operations of committed transactions from main memory buffers to disk,
  - Write a checkpoint record into the log file and write log to disk
  - Resume executing transactions
- Only changes made by transactions committed between the last checkpoint and a system failure have to be redone

# Redo With Checkpoint



Changes made by  $T_1$  are stored in the database, so there is nothing to do with them

DBMS will **REDO** transactions  $T_2$  and  $T_3$  (their logs are already on disk)

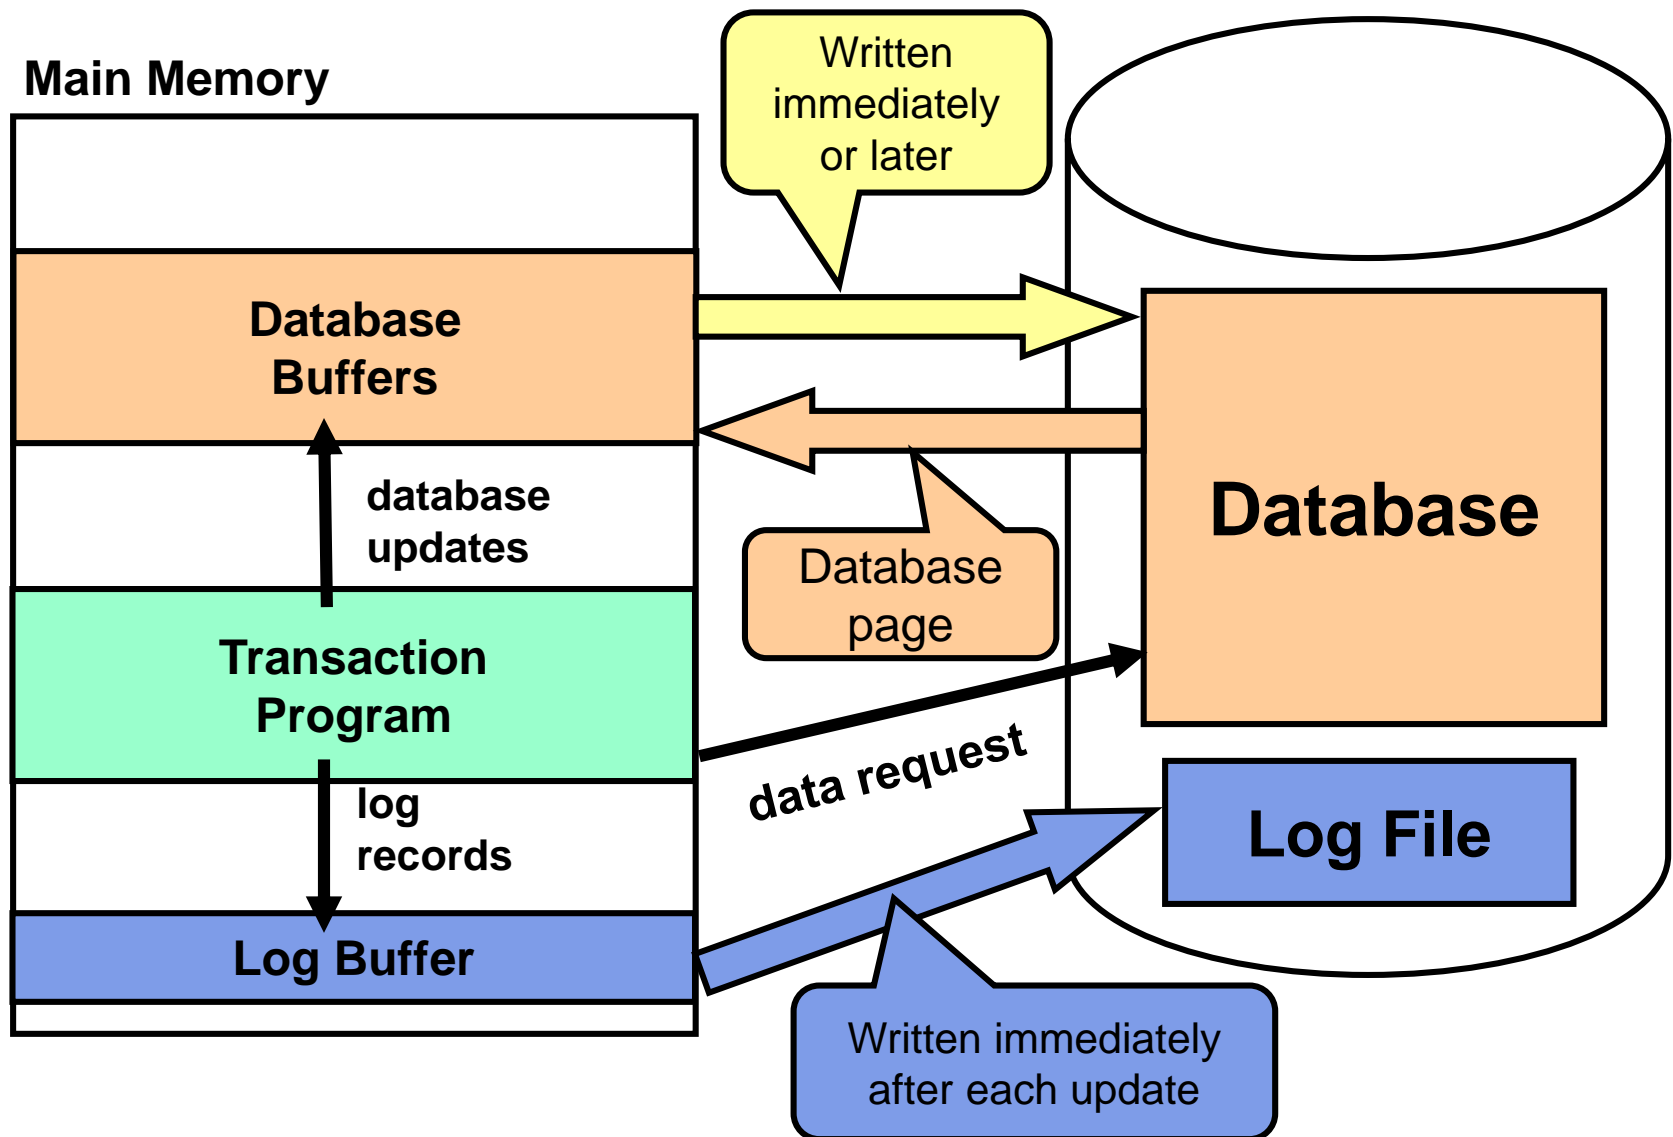
DBMS or the user has to **rerun**  $T_4$  and  $T_5$



# Immediate Update

- The main idea:
  - When a transaction issues an update command, before and after images are recorded into a log file on **disk**, and (thereupon) the database **can** (but do not have to) be **immediately** updated
- There are two versions of immediate update algorithm
  - One writes all updates from buffers into the database **before** the transaction actually **commits**
  - The other one allows a transaction to **commit before** all its updates have been written in the database
  - The second version is the most common one

# Immediate Update Layout




# Transaction Roll-Back (**Undo**)

- If a transaction fails for any reason, its effects are rolled-back
- **Roll-back** of a transaction is a procedure of **undoing** changes made against the database by a non-committed transaction
- A roll-back is done by restoring the **before images** (old values) of each database item changed by the transaction
- The roll-back is done in the **reverse order** of the order the operations were written into the log file on disk

# Redo in Immediate Update

- In the case of a system crash:
  - Because there can be some updates of the committed transactions that are not written into the database, effects of the committed transactions are redone to the database
    - Recall, we allow a transaction to commit before all of its updates have been written into the database
- Only the **last** updates of each item have to be redone

# Immediate Update



Changes made by  $T_1$  are stored in the database, so there is nothing to do with them

DBMS will **REDO** transactions  $T_2$  and  $T_3$  (their logs are already on disk)

DBMS has to **UNDO**  $T_4$  and  $T_5$

# A Question for You

- Immediate update with roll-back requires redoing transactions that committed between the last checkpoint and the moment of system crash
- What if the changes are already made against the database?
  - a) Redoing will bring the database into an inconsistent state
  - b) This question is too complex, and can not be answered without a careful analysis
  - c) Redoing will make no harm to the database, because it will bring the database in the same consistent state as it was in

# Summary

- *Recovery* from a *catastrophic* failure is made by applying operations of committed transactions from archived log files on an archived database back up copy
- Recovery from a *noncatastrophic* failure can be accomplished through:
  - Deferred update,
  - Immediate update, and
  - Other recovery schemes
- *Deferred update* means writing changes into a database after a transaction commits, but the log is written to disk just before the transaction actually commits
- *Immediate update* means that changes are immediately stored in a log file on disk. So, changes can be written into a database before or after a transaction commits