

# Structured Query Language (SQL) Tutorial

SWEN304/SWEN439

Lecturer: Dr Hui Ma

**Engineering and Computer Science**



# Outline

- SQL Constraints:
  - CHECK constraint
  - Referential integrity constraint  
MATCH PARTIAL | MATCH FULL | MATCH SIMPLE
- Using the same table in different context
- Correlated queries

# Specifying CHECK Constraint

```
CREATE TABLE COURSE (  
  CourId CHAR(4) CONSTRAINT cspk PRIMARY KEY
```

CHECK constraint comes here

```
,  
  CName CHAR(15) NOT NULL,  
  Points INT NOT NULL,  
  Dept CHAR(25) );
```

- Suppose we would like to define an additional constraint on CourId that restricts this attribute values to those that follow the pattern:
  - the first character is a capital letter,
  - the next three characters are numbers between 100 and 999

# Additional Constraint on CourId

```
CREATE TABLE COURSE (
  CourId CHAR(4) CONSTRAINT cspk PRIMARY KEY
  CONSTRAINT valcon
    CHECK ( (SUBSTR(CourId,1,1) BETWEEN 'A' AND 'Z')
    AND (SUBSTR(CourId,2,3) SIMILAR TO '[1-9][0-9][0-9]')),
  CName CHAR(15) NOT NULL,
  Points INT NOT NULL,
  Dept CHAR(25) );
```

- Here, we used **regular expressions** to define a CHECK constraint of the *format (pattern)* type

# Referential Integrity – A Formal Definition

- Relations  $r(N_2)$  and  $r(N_1)$  satisfy referential integrity  $N_2[Y] \subseteq N_1[X]$  if

$$(\forall u \in r(N_2))(\exists v \in r(N_1))(u[Y] = v[X] \vee (\exists i \in \{1, \dots, m\})(u[B_i] = \omega))$$

- Either tuples  $u$  and  $v$  are equal on  $X$  and  $Y$  values, or there exists at least one attribute in  $Y$  whose  $u$  value is null

# Referential Integrity (1)

- TEXT\_BOOK ({Title, ISBN, C\_Code, C\_Num}, {ISBN })
- COURSE ({C\_Code, C\_Num, C\_Name },  
          {C\_Code + C\_Num,})
- How do we specify the referential integrity constraint:  
    TEXT\_BOOK [C\_Code, C\_Num]  $\subseteq$   
    COURSE [C\_Code, C\_Num]

Need:

- **Referring** and **Referred** relational variables and fields
- No Match clause or Match: **FULL | PARTIAL | SIMPLE**
- Action: **NO ACTION, CASCADE, SET NULL, SET DEFAULT**

# Referential Integrity (2) MATCH

- MATCH clause:

SIMPLE /no MATCH clause specified (Default):

- Two tuples either match on primary key/foreign key values, or the foreign key has at least one null valued component

PARTIAL:

- Two tuples either match on primary key/foreign key values, or the not null valued foreign key components match the corresponding components of at least one primary key value

FULL:

- Two tuples either match on primary key/foreign key values, or all foreign key components are null

# Referential Integrity (3)

```
CREATE TABLE COURSE (
    C_Code CHAR(4) NOT NULL DEFAULT 'stat',
    C_Num INT CHECK(C_NUM BETWEEN 100 AND 999) DEFAULT 100,
    C_Name CHAR(25),
    PRIMARY KEY (C_Code, C_Num )
);
```

```
CREATE TABLE TEXT_BOOK (
    Title CHAR(30) NOT NULL,
    ISBN INT PRIMARY KEY,
    C_Code CHAR(4),
    C_Num INT(2) CHECK(C_Num > 99 AND C_Num < 1000),
    FOREIGN KEY (C_Code, C_Num) REFERENCES COURSE
        [MATCH <condition>] ON DELETE <action>
);
```



# Referential Integrity (4a)

**[MATCH <condition>] ON DELETE <action>**

- **No MATCH clause** (Default)
- **<action>: NO ACTION** (RESTRICT)

**TEXTBOOK**

Title	ISBN	C_Code	C_Num
FDBS	1111	comp	302

**COURSE**

C_Code	C_Num	Name
comp	302	DB

INSERT INTO TEXTBOOK VALUES ('Disc. Log.', 2222, 'math', null);

??

DELETE FROM COURSE WHERE C\_Code = 'comp' AND C\_Num = 302;

??

# Referential Integrity (4b - Answer)

**[MATCH <condition>] ON DELETE <action>**

- **No MATCH clause** (Default)
- **<action>: NO ACTION** (RESTRICT)

**TEXTBOOK**

Title	ISBN	C_Code	C_Num
FDBS	1111	comp	302

**COURSE**

C_Code	C_Num	Name
comp	302	DB

INSERT INTO TEXTBOOK VALUES ('Disc. Log.', 2222, 'math', null);  
*Successful, because of MATCH default*

DELETE FROM COURSE WHERE C\_Code = 'comp' AND C\_Num = 302;  
*Rejected, because of NO ACTION*

# Referential Integrity (5a)

**[MATCH <condition>] ON DELETE <action>**

- **MATCH <condition>: PARTIAL**
- **<action>: CASCADE,**

**TEXTBOOK**

Title	<u>ISBN</u>	C_Code	C_Num
FDBS	1111	comp	302

**COURSE**

<u>C_Code</u>	<u>C_Num</u>	Name
comp	302	DB

```
INSERT INTO TEXTBOOK VALUES ('Disc. Log.', 2222, 'math', null);
??
```

```
INSERT INTO TEXTBOOK VALUES ('Disc. Log.', 2222, 'comp', null);
??
```

```
DELETE FROM COURSE WHERE C_Code = 'comp' AND C_Num = 302;
??
```

# Referential Integrity (5b - Answer)

**[MATCH <condition>] ON DELETE <action>**

- **MATCH <condition>: PARTIAL**
- **<action>: CASCADE,**

**TEXTBOOK**

Title	<u>ISBN</u>	C_Code	C_Num
FDBS	1111	comp	302

**COURSE**

<u>C_Code</u>	<u>C_Num</u>	Name
comp	302	DB

INSERT INTO TEXTBOOK VALUES ('Disc. Log.', 2222, 'math', null);  
*Rejected, because of MATCH PARTIAL*

INSERT INTO TEXTBOOK VALUES ('Disc. Log.', 2222, 'comp', null);  
*Successful, because of MATCH PARTIAL*

DELETE FROM COURSE WHERE C\_Code = 'comp' AND C\_Num = 302;  
*Successful, because of CASCADE (all tuples will be deleted)*

# Referential Integrity (6a)

**[MATCH <condition>] ON DELETE <action>**

- **MATCH <condition>: FULL**
- **<action>: SET NULL,**

**TEXTBOOK**

Title	<u>ISBN</u>	C_Code	C_Num
FDBS	1111	comp	302

**COURSE**

<u>C_Code</u>	<u>C_Num</u>	Name
comp	302	DB

```
INSERT INTO TEXTBOOK VALUES ('Disc. Log.', 2222, comp, null);
??
```

```
INSERT INTO TEXTBOOK VALUES ('Disc. Log.', 2222, null, null);
?
```

```
DELETE FROM COURSE WHERE C_Code = 'comp' AND C_Num = 302;
??
```

# Referential Integrity (6b - Answer)

**[MATCH <condition>] ON DELETE <action>**

- **MATCH <condition>: FULL**
- **<action>: SET NULL,**

**TEXTBOOK**

Title	<u>ISBN</u>	C_Code	C_Num
FDBS	1111	comp	302

**COURSE**

<u>C_Code</u>	<u>C_Num</u>	Name
comp	302	DB

**INSERT INTO TEXTBOOK VALUES ('Disc. Log.', 2222, 'comp', null);**

*Rejected, because of MATCH FULL*

**INSERT INTO TEXTBOOK VALUES ('Disc. Log.', 2222, null, null);**

*Successful, because of MATCH FULL*

**DELETE FROM COURSE WHERE C\_Code = comp AND C\_Num = 302;**

*Successful, because of SET NULL (course tuple will be deleted, and the foreign key of the textbook tuple will be nullified)*

# Referential Integrity (7a)

**[MATCH <condition>] ON DELETE <action>**

- **<action>: SET DEFAULT,**

**TEXTBOOK**

Title	<u>ISBN</u>	C_Code	C_Num
FDBS	1111	comp	302

**COURSE**

<u>C_Code</u>	<u>C_Num</u>	Name
comp	302	DB
stat	100	STAT

DELETE FROM *COURSE* WHERE *C\_Code* = 'comp' AND *C\_Num* = 302;

??

# Referential Integrity (7b - Answer)

**[MATCH <condition>] ON DELETE <action>**

- **<action>: SET DEFAULT**

**TEXTBOOK**

Title	<u>ISBN</u>	C_Code	C_Num
FDBS	1111	comp	302

**COURSE**

<u>C_Code</u>	<u>C_Num</u>	Name
comp	302	DB
stat	100	STAT

**DELETE FROM COURSE WHERE C\_Code = 'comp' AND C\_Num = 302;**

*Successful, because of SET DEFAULT (course tuple will be deleted, and the foreign key of the **textbook tuple** will be set to *stat 100*)*



## STUDENT

LName	FName	StudId	Major	TrPts
Smith	Susan	131313	Comp	54
Bond	James	007007	Math	120
Smith	Susan	555555	Comp	30
Cecil	John	010101	Math	90

## COURSE

CName	CrsId	Points	Dept
DB Sys	C302	15	Comp
SofEng	C301	15	Comp
DisMat	M214	22	Math
Pr&Sys	C201	22	Comp

## GRADES

StudId	CrsId	Grade
007007	C302	A+
555555	C302	ω
007007	C301	A
007007	M214	A+
131313	C201	B-
555555	C201	C
131313	C302	ω
007007	C201	A
010101	C201	ω

# A Question for You (Tricky Null Value)

- What is wrong with the following query:

```
SELECT *
FROM GRADES
WHERE Grade = Null;
```

since it returns an empty table

StudId	CrsId	Grade

- Answer:
  - There is a mistake, only I do not know where
  - PostgreSQL is rubbish
  - Null is not a real value. It can be anything. So, to the questions whether

Null = Null, or

Grade = Null

PostgreSQL answers "I don't know".

# Multiple uses of the same table

- SQL allows multiple occurrences of the **same** table in a **FROM** clause
- In that case, each occurrence of the same table has a **different role**, or a different context of usage
- Aliases are used to denote the context of usage

# Multiple Uses of the Same Table

- **Query:** Retrieve student *ids* and *TrPts* of students that have greater number of transfer points than the student with *StudentId* = 131313

```
SELECT s1.StudId, s1.TrPts
FROM STUDENT s1, STUDENT s2
WHERE s1.TrPts > s2.TrPts    AND
      s2.StudId = 131313 ;
```

Not an  
Equi Join

- The context of *s2* is "number of points of the student with *StudentId* = 131313"
- The context of *s1* is "list of students having greater number of points than student with *StudentId* = 131313"

# Nested Queries

- Some queries require **comparing** a tuple to a collection of tuples (e.g., *students doing courses that have more than 100 students*)
- This task can be accomplished by **embedding** a SQL query into `WHERE` clause of another query
  - The embedded query is called **nested query**,
  - The query containing the nested query is called **outer query**
- The **comparison** is made by using `IN`,  $\theta$  `ANY`,  $\theta$  `SOME`, and  $\theta$  `ALL` operators, where  $\theta \in \{ =, <, <=, >=, >, < > \}$
- **Note:** `IN`  $\Leftrightarrow$  `=ANY`    and    `IN`  $\Leftrightarrow$  `=SOME`

# Correlated Nested Queries

- Let the variable  $s$  contain the current tuple of the outer query
- If the nested query doesn't refer to  $s$  :
  - The nested query computes the same result for each tuple in  $s$
  - The outer query and the nested query are said to be **uncorrelated**
- If a condition in the `WHERE` clause of the nested query refers to some attributes of a relation declared in the outer query, the two queries are said to be **correlated**
  - Have to compute the inner query for **each** tuple considered by the outer query
  - Correlated nested queries consume **more** computation time than uncorrelated ones

# Nested Queries

- Is the following nested query correlated or not?
  - Select first names of the students that didn't enroll M214

```
SELECT s1.FName
FROM STUDENT s1
WHERE s1.StudId IN
    ((SELECT s2.StudId FROM STUDENT s2)
    EXCEPT
    (SELECT StudId FROM GRADES
    WHERE CourId = 'M214'));
```

FName
Susan
Susan
John

# Correlated Nested Queries

- Consider the relation schemas:

EMPLOYEE ({EmpId, EmpName, Salary }, {EmpId } )

PROJECT ({ProjId, ProjName }, {ProjId })

WORKS\_ON ({EmpId, ProjId, NoOfHours}, {EmpId + ProjId })



# Correlated Nested Queries (1)

- **Query:** Retrieve names of employees that work more hours on a project than the average number of hours on that same project :

```
SELECT e.EmpName
FROM EMPLOYEE e, WORKS_ON w
WHERE e.EmpId = w.EmpId AND
      NoOfHours > (SELECT AVG(NoOfHours )
                   FROM WORKS_ON w1
                   WHERE w.ProjId = w1.ProjId );
```

- Here, in the nested query, the task of  $w$  is to focus on the current project, and allow computing requested average
- $w.ProjId$  is a **correlated** attribute

## Correlated Nested Queries (2)

- Query: Retrieve names of employees that are ranked first three according to their salaries (each employee has a different salary)

```
SELECT e.EmpName
FROM EMPLOYEE e
WHERE 3 > (
    SELECT COUNT(*)
    FROM EMPLOYEE e1
    WHERE e1.Salary > e.Salary);
```

- In the nested query, only the employees that have higher salary than the current employee are selected, and current employee will be selected in the outer query only if the number of employees, selected in the inner query, is less than 3
- e.Salary is a **correlated** attribute

## Another Query

- Query: show the project name and the average salary of employees who worked on projects that took a total of more than 1000 hours.

```
CREATE VIEW ExpensiveProjects AS
  (SELECT ProjName, AVG(Salary)
   FROM WORKS_ON NATURAL JOIN EMPLOYEE NATURAL
   JOIN PROJECT
   WHERE Projid IN
    (SELECT ProjId
     FROM (SELECT ProjId, SUM(NoOfHours) AS TotHours
           FROM WORKS_ON
           GROUP BY ProjId) AS ProjHours
     WHERE TotHours > 1000)
   GROUP BY ProjName
  );
```

## Another Query

- Query: show the project name and the average salary of employees who worked on projects that took a total of more than 1000 hours.

```
CREATE VIEW ExpensiveProjects AS
  (SELECT ProjName, AVG(Salary)
   FROM WORKS_ON NATURAL JOIN EMPLOYEE NATURAL JOIN
   (SELECT ProjId, SUM(NoOfHours) AS TotHours
    FROM WORKS_ON
   GROUP BY ProjId) AS ProjHours
   WHERE TotHours > 1000
  GROUP BY ProjName
 );
```

# Exercises

## COMPANY Database Schema

### EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

### DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

### DEPT\_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

### PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

### WORKS\_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

### DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------

**Figure 5.5**  
Schema diagram for  
the COMPANY  
relational database  
schema.

# Exercises

In SQL, specify the following queries on the COMPANY database using the concept of nested queries.

- 1) Retrieve the names of all employees who work in the department that has the employee with the highest salary among all employees.
- 2) Retrieve the names of all employees whose supervisor's supervisor has '888665555' for Ssn.
- 3) Retrieve the names of employees who make at least \$10,000 more than the employee who is paid the least in the company.

# Exercises

- 1) Retrieve the names of all employees who work in the department that has the employee with the highest salary among all employees.

```
SELECT LNAME
FROM EMPLOYEE
WHERE DNO = (SELECT DNO
              FROM EMPLOYEE
              WHERE SALARY = (SELECT MAX(SALARY)
                             FROM EMPLOYEE) )
```

# Exercises

- 2) Retrieve the names of all employees whose supervisor's supervisor has '888665555' for Ssn.

```
SELECT LNAME  
FROM EMPLOYEE  
WHERE SUPERSSN IN  
    (SELECT SSN  
     FROM EMPLOYEE  
     WHERE SUPERSSN = '888665555' )
```



# Exercises

- 3) Retrieve the names of employees who make at least \$10,000 more than the employee who is paid the least in the company

```
SELECT LNAME  
FROM EMPLOYEE  
WHERE SALARY >= 10000 +  
      ( SELECT MIN(SALARY)  
        FROM EMPLOYEE)
```